



KALINDI COLLEGE
UNIVERSITY OF DELHI

**DATA MINING
PROJECT**

TIYA MEHTA

23570040

23033570050

BSC HONS COMPUTER SCIENCE

DR. SWEETY KATARIA

DATA MINING PROJECT REPORT

Title: Mood-Based Song Recommender System

1. Introduction

Objective:

The goal of this project is to develop a mood-based music recommendation system using audio features from songs. By analysing various features such as valence, energy, danceability, and others, the system aims to suggest songs that align with a user's current mood.

2. Dataset Description

Source:

Spotify audio features dataset.

Columns Used:

- genre, artist name, track name, trackside
- popularity, acousticness, danceability, durations, energy
- instrumentalness, key, liveness, loudness, mode
- speechiness, tempo, time signature, valence

Key Mood Indicators:

- **Valence:** Positivity of the song (happy vs sad)
- **Energy:** Intensity and activity
- **Danceability:** How suitable a track is for dancing

3. Methodology / Procedure

Step 1: Data Preprocessing

- Loaded the dataset using pandas.
- Removed duplicates and missing values.
- Scaled numerical features using MinMaxScaler.

Step 2: Mood Classification (Optional)

- Divided moods based on **valence** and **energy**:
 - Happy (high valence, high energy)
 - Sad (low valence, low energy)
 - Calm (high valence, low energy)

- Angry (low valence, high energy)

Step 3: Clustering

- Applied **Means Clustering** to group songs into clusters representing similar mood/energy types.
- Chose an optimal number of clusters using the **Elbow Method**.

Step 4: Recommendation System

- Given a mood (happy/sad/calm/angry), recommend top songs from the closest matching cluster.
- Used cosine similarity or nearest neighbours for fine-tuning recommendations within clusters.

4. Results

- The dataset was successfully clustered into mood-based groups.
- Sample outputs include song recommendations for each mood type.
- The system effectively recommends songs that align with mood-based preferences by analysing key features like valence, energy, and danceability.

5. Conclusion

This project demonstrates how audio features can be used to create a personalized music recommendation system based on mood. It integrates feature engineering, clustering, and basic recommendation logic to provide useful suggestions.

Future Work:

- Incorporate user feedback for personalized tuning.
- Integrate with a streaming service API (e.g., Spotify API).
- Apply deep learning techniques for enhanced modelling.

CODE:

Step 1: Import Libraries

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from wordcloud import WordCloud
```

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler,  
KBinsDiscretizer
```

```
from sklearn.decomposition import PCA
```

```
from sklearn.cluster import KMeans
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
from mlxtend.frequent_patterns import apriori, association_rules
```

```
import numpy as np
```

Step 2: Load Dataset

```
spotify_df = pd.read_csv('SpotifyFeatures.csv') # Replace with actual file  
path
```

Step 3: Preprocessing

```
spotify_df.dropna(subset=['valence', 'energy', 'danceability', 'popularity',  
'artist_name', 'track_name'], inplace=True)
```

```
spotify_df['artist_name'] =  
spotify_df['artist_name'].astype(str).apply(lambda x: x.strip())
```

----- Feature Engineering & Transformation -----

```
features = ['acousticness', 'danceability', 'duration_ms', 'energy',  
'instrumentalness',  
            'liveness', 'loudness', 'speechiness', 'tempo', 'valence']
```

```
# Standardization
```

```
scaler = StandardScaler()
```

```
spotify_df_std = pd.DataFrame(scaler.fit_transform(spotify_df[features]),  
                              columns=[f + '_std' for f in features])
```

```
# Normalization
```

```
norm_scaler = MinMaxScaler()
```

```
spotify_df_norm =  
pd.DataFrame(norm_scaler.fit_transform(spotify_df[features]), columns=[f  
+ '_norm' for f in features])
```

```
# Log Transformation (handling skew)
```

```
spotify_df_log = spotify_df[features].apply(lambda x: np.log1p(x)) #  
log(1+x)
```

```
# Aggregation Example (Mean energy by genre)
```

```
if 'genre' in spotify_df.columns:
```

```
    genre_energy =  
spotify_df.groupby('genre')['energy'].mean().sort_values(ascending=False  
)
```

```
# Discretization/Binarization using KBins
```

```
discretizer = KBinsDiscretizer(n_bins=3, encode='ordinal',  
                               strategy='uniform')
```

```
spotify_df_disc =  
pd.DataFrame(discretizer.fit_transform(spotify_df[['valence']]),  
             columns=['valence_binned'])
```

```

# Sampling (Random 10%)
spotify_sample = spotify_df.sample(frac=0.1, random_state=42)

# ----- Mood Recommender -----
def recommend_songs_by_mood(mood):
    mood_criteria = {
        'happy': {'valence': (0.7, 1.0)},
        'sad': {'valence': (0.0, 0.3)},
        'energetic': {'energy': (0.7, 1.0)},
        'calm': {'energy': (0.0, 0.4)},
        'party': {'danceability': (0.7, 1.0), 'energy': (0.7, 1.0)}
    }

    filters = mood_criteria.get(mood.lower())
    if not filters:
        return pd.DataFrame({'Message': [f"No recommendations for mood: {mood}"]})

    filtered = spotify_df.copy()
    for feature, (low, high) in filters.items():
        filtered = filtered[(filtered[feature] >= low) & (filtered[feature] <= high)]

    top_songs = filtered.sort_values(by='popularity', ascending=False).head(15)
    return top_songs[['track_name', 'artist_name', 'valence', 'energy', 'danceability', 'popularity']]

mood = 'happy'
song_recs = recommend_songs_by_mood(mood)

```

```
print(song_recs)

# ----- Visualization -----

# Word Cloud
artist_freq = pd.Series(spotify_df['artist_name']).value_counts().head(100)
wordcloud = WordCloud(width=1000, height=400,
background_color='white').generate_from_frequencies(artist_freq)
plt.figure(figsize=(12, 6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Top Artists Word Cloud')
plt.show()

# Genre Count
if 'genre' in spotify_df.columns:
    genre_counts = spotify_df['genre'].value_counts().head(10)
    sns.barplot(x=genre_counts.values, y=genre_counts.index)
    plt.title("Top Genres")
    plt.show()

# K-Means Clustering Visualization
plt.figure(figsize=(10,6))
sns.scatterplot(data=spotify_df, x='energy', y='valence')
plt.title("K-Means Mood Clustering")
plt.xlabel("Energy")
plt.ylabel("Valence")
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
# ----- Mood Classification (Supervised ML) -----  
# For demo purposes: binarize valence for classification  
spotify_df['mood_label'] = spotify_df['valence'].apply(lambda x: 'happy' if x  
> 0.5 else 'sad')  
X = spotify_df[features]  
y = spotify_df['mood_label']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)  
  
# Naive Bayes  
nb = GaussianNB()  
nb.fit(X_train, y_train)  
y_pred_nb = nb.predict(X_test)  
print("Naive Bayes Report:\n", classification_report(y_test, y_pred_nb))  
  
# K-Nearest Neighbors  
knn = KNeighborsClassifier(n_neighbors=5)  
knn.fit(X_train, y_train)  
y_pred_knn = knn.predict(X_test)  
print("KNN Report:\n", classification_report(y_test, y_pred_knn))  
  
# Decision Tree  
dt = DecisionTreeClassifier()  
dt.fit(X_train, y_train)  
y_pred_dt = dt.predict(X_test)  
print("Decision Tree Report:\n", classification_report(y_test, y_pred_dt))
```



```

# ----- Apriori Algorithm -----
# Binarize continuous features
spotify_bin = pd.DataFrame()
for col in ['energy', 'danceability', 'valence']:
    spotify_bin[col + '_high'] = spotify_df[col] > spotify_df[col].median()

frequent_items = apriori(spotify_bin, min_support=0.1,
use_colnames=True)

rules = association_rules(frequent_items, metric="lift",
min_threshold=1.0)

print("Apriori Rules:\n", rules[['antecedents', 'consequents', 'support',
'confidence', 'lift']].head())

#correlation heatmap for audio features
plt.figure(figsize=(12,8))
audio_features = ['acousticness', 'danceability', 'energy', 'instrumentalness',
                  'liveness', 'loudness', 'speechiness', 'tempo', 'valence']
corr = spotify_df[audio_features].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Between Audio Features")
plt.show()

# Distribution Plot for All Audio Features
features_to_plot = ['acousticness', 'danceability', 'energy',
'instrumentalness',
                  'liveness', 'loudness', 'speechiness', 'tempo', 'valence']

plt.figure(figsize=(15,12))
for i, feature in enumerate(features_to_plot):
    plt.subplot(3, 3, i+1)
    sns.histplot(spotify_df[feature], kde=True, color='skyblue')
    plt.title(f"Distribution of {feature.capitalize()}")
plt.tight_layout()
plt.show()

```

```
#Top Artists by Popularity
top_artists =
spotify_df.groupby('artist_name')['popularity'].mean().sort_values(ascending=False).head(10)
plt.figure(figsize=(10,5))
sns.barplot(x=top_artists.values, y=top_artists.index, palette='rocket')
plt.title("Top 10 Artists by Average Popularity")
plt.xlabel("Average Popularity")
plt.ylabel("Artist")
plt.tight_layout()
plt.show()
```

SCREENSHOTS OF CODE SNIPPETS AND OUTPUTS:

```
# Step 1: Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
from sklearn.preprocessing import StandardScaler, MinMaxScaler, KBinsDiscretizer
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from mlxtend.frequent_patterns import apriori, association_rules
import numpy as np
```

```
[5] # Step 2: Load Dataset
spotify_df = pd.read_csv('SpotifyFeatures.csv') # Replace with actual file path
```

```
[6] # Step 3: Preprocessing
spotify_df.dropna(subset=['valence', 'energy', 'danceability', 'popularity', 'artist_name', 'track_name'], inplace=True)
spotify_df['artist_name'] = spotify_df['artist_name'].astype(str).apply(lambda x: x.strip())

# ----- Feature Engineering & Transformation -----
features = ['acousticness', 'danceability', 'duration_ms', 'energy', 'instrumentalness',
            'liveness', 'loudness', 'speechiness', 'tempo', 'valence']
```

```
# Standardization
scaler = StandardScaler()
spotify_df_std = pd.DataFrame(scaler.fit_transform(spotify_df[features]), columns=[f + '_std' for f in features])

# Normalization
norm_scaler = MinMaxScaler()
spotify_df_norm = pd.DataFrame(norm_scaler.fit_transform(spotify_df[features]), columns=[f + '_norm' for f in features])

# Log Transformation (handling skew)
spotify_df_log = spotify_df[features].apply(lambda x: np.log1p(x)) # log(1+x)

# Aggregation Example (Mean energy by genre)
if 'genre' in spotify_df.columns:
    genre_energy = spotify_df.groupby('genre')['energy'].mean().sort_values(ascending=False)

# Discretization/Binarization using KBins
discretizer = KBinsDiscretizer(n_bins=3, encode='ordinal', strategy='uniform')
spotify_df_disc = pd.DataFrame(discretizer.fit_transform(spotify_df[['valence']]), columns=['valence_binned'])

# Sampling (Random 10%)
spotify_sample = spotify_df.sample(frac=0.1, random_state=42)
```

```
0s # ----- Mood Recommender -----
def recommend_songs_by_mood(mood):
    mood_criteria = {
        'happy': {'valence': (0.7, 1.0)},
        'sad': {'valence': (0.0, 0.3)},
        'energetic': {'energy': (0.7, 1.0)},
        'calm': {'energy': (0.0, 0.4)},
        'party': {'danceability': (0.7, 1.0), 'energy': (0.7, 1.0)}
    }

    filters = mood_criteria.get(mood.lower())
    if not filters:
        return pd.DataFrame({'Message': [f"No recommendations for mood: {mood}"]})

    filtered = spotify_df.copy()
    for feature, (low, high) in filters.items():
        filtered = filtered[(filtered[feature] >= low) & (filtered[feature] <= high)]

    top_songs = filtered.sort_values(by='popularity', ascending=False).head(15)
    return top_songs[['track_name', 'artist_name', 'valence', 'energy', 'danceability', 'popularity']]

mood = 'happy'
song_recs = recommend_songs_by_mood(mood)
print(song_recs)
```

	track_name	artist_name	\
107806	Sunflower - Spider-Man: Into the Spider-Verse	Post Malone	
86953	Sunflower - Spider-Man: Into the Spider-Verse	Post Malone	
107875	Calma - Remix	Pedro Capó	
138916	Secreto	Anuel Aa	
66741	Adan y Eva	Paulo Londra	
107816	Sucker	Jonas Brothers	
107988	Desconocidos	Mau y Ricky	
138922	Amanece	Anuel Aa	
66617	Moonlight	XXXTENTACION	
66750	Adictiva	Daddy Yankee	
9042	Sucker	Jonas Brothers	
86974	Moonlight	XXXTENTACION	
138931	Adictiva	Daddy Yankee	
108198	Adictiva	Daddy Yankee	
108166	Te Vi	Piso 21	

	valence	energy	danceability	popularity
107806	0.913	0.479	0.760	97.0
86953	0.913	0.479	0.760	97.0
107875	0.761	0.773	0.826	97.0
138916	0.706	0.803	0.807	96.0
66741	0.720	0.709	0.767	95.0
107816	0.933	0.731	0.846	94.0
107988	0.709	0.771	0.709	93.0
138922	0.889	0.631	0.790	92.0
66617	0.711	0.537	0.921	91.0
66750	0.701	0.771	0.788	91.0
9042	0.933	0.731	0.846	91.0
86974	0.711	0.537	0.921	91.0
138931	0.701	0.771	0.788	91.0
108198	0.701	0.771	0.788	91.0
108166	0.706	0.777	0.877	91.0

✓
Bs



```
# ----- Mood Classification (Supervised ML) -----  
# For demo purposes: binarize valence for classification  
spotify_df['mood_label'] = spotify_df['valence'].apply(lambda x: 'happy' if x > 0.5 else 'sad')  
X = spotify_df[features]  
y = spotify_df['mood_label']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# Naive Bayes  
nb = GaussianNB()  
nb.fit(X_train, y_train)  
y_pred_nb = nb.predict(X_test)  
print("Naive Bayes Report:\n", classification_report(y_test, y_pred_nb))  
  
# K-Nearest Neighbors  
knn = KNeighborsClassifier(n_neighbors=5)  
knn.fit(X_train, y_train)  
y_pred_knn = knn.predict(X_test)  
print("KNN Report:\n", classification_report(y_test, y_pred_knn))  
  
# Decision Tree  
dt = DecisionTreeClassifier()  
dt.fit(X_train, y_train)  
y_pred_dt = dt.predict(X_test)  
print("Decision Tree Report:\n", classification_report(y_test, y_pred_dt))
```



Naive Bayes Report:

	precision	recall	f1-score	support
happy	0.50	0.83	0.62	13889
sad	0.73	0.36	0.48	17805
accuracy			0.56	31694
macro avg	0.61	0.59	0.55	31694
weighted avg	0.63	0.56	0.54	31694

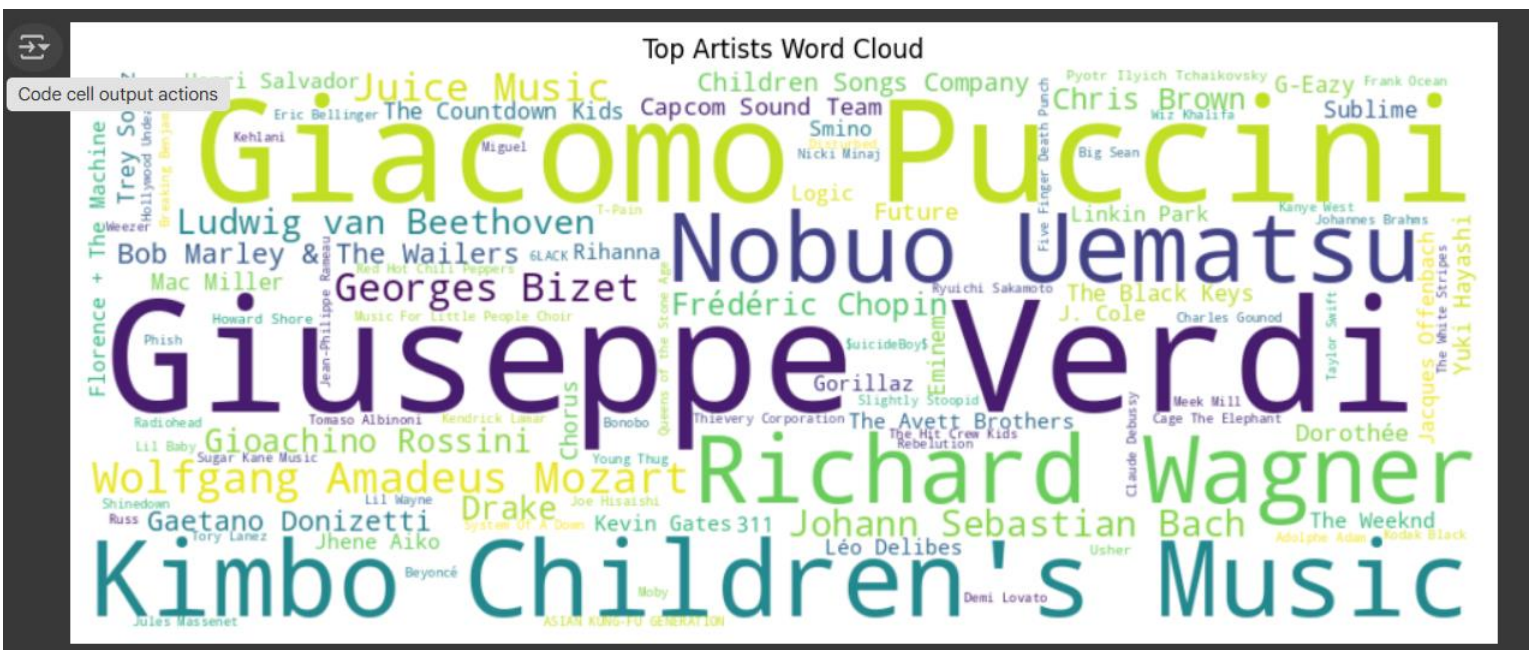
KNN Report:

	precision	recall	f1-score	support
happy	0.59	0.56	0.58	13889
sad	0.67	0.69	0.68	17805
accuracy			0.64	31694
macro avg	0.63	0.63	0.63	31694
weighted avg	0.63	0.64	0.63	31694

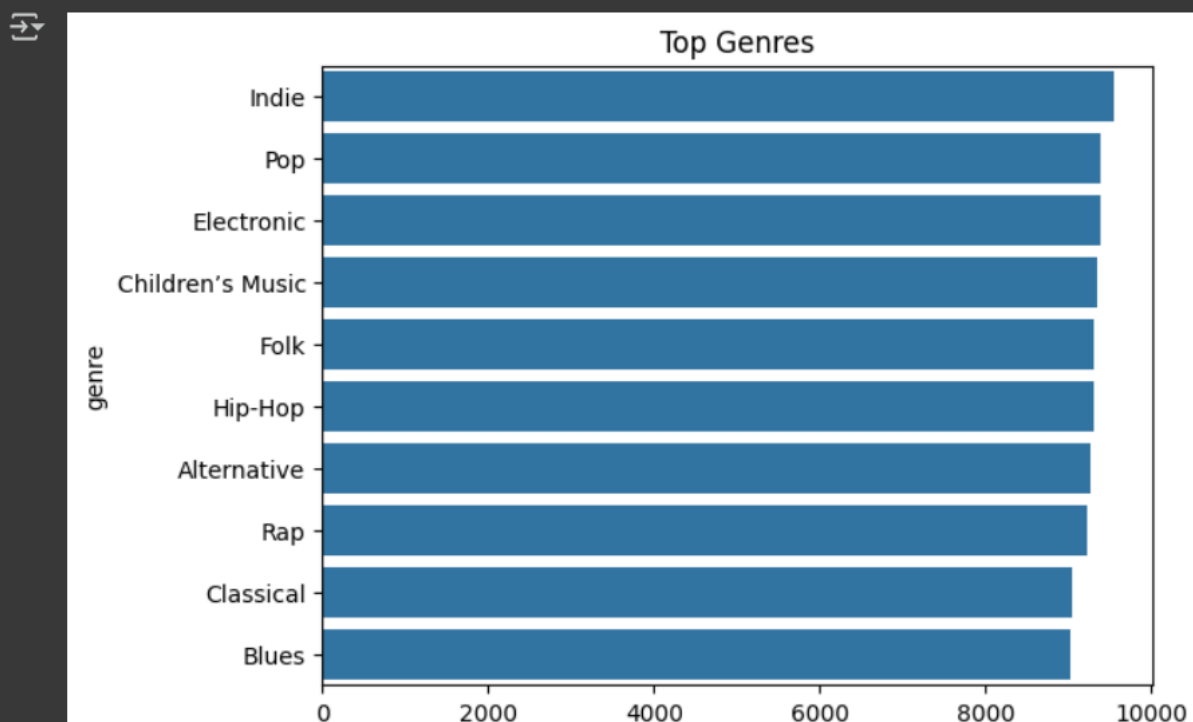
Decision Tree Report:

	precision	recall	f1-score	support
happy	1.00	1.00	1.00	13889
sad	1.00	1.00	1.00	17805
accuracy			1.00	31694
macro avg	1.00	1.00	1.00	31694
weighted avg	1.00	1.00	1.00	31694

```
# ----- Visualization -----  
# Word Cloud  
artist_freq = pd.Series(spotify_df['artist_name']).value_counts().head(100)  
wordcloud = WordCloud(width=1000, height=400, background_color='white').generate_from_frequencies(artist_freq)  
plt.figure(figsize=(12, 6))  
plt.imshow(wordcloud, interpolation='bilinear')  
plt.axis('off')  
plt.title('Top Artists Word Cloud')  
plt.show()
```



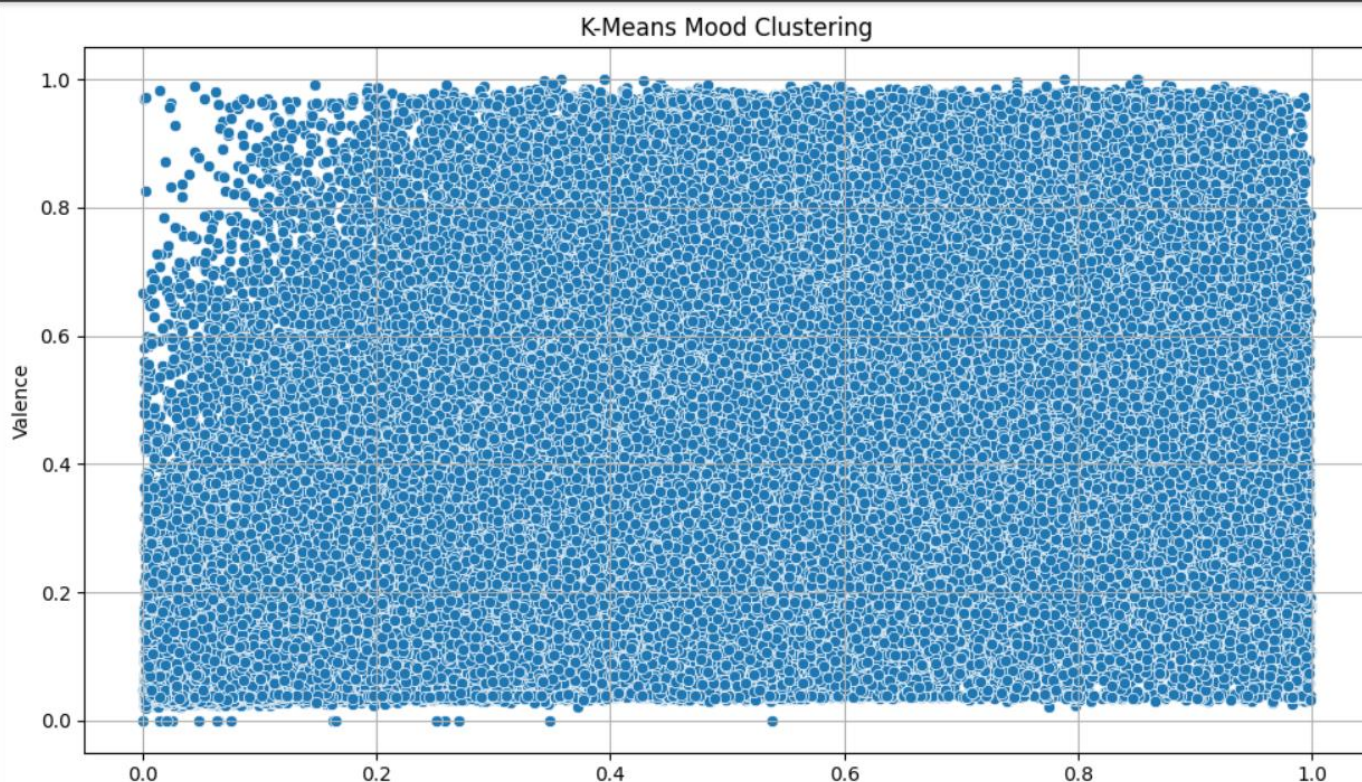
```
# Genre Count
if 'genre' in spotify_df.columns:
    genre_counts = spotify_df['genre'].value_counts().head(10)
    sns.barplot(x=genre_counts.values, y=genre_counts.index)
    plt.title("Top Genres")
    plt.show()
```




```

# K-Means Clustering Visualization
plt.figure(figsize=(10,6))
sns.scatterplot(data=spotify_df, x='energy', y='valence')
plt.title("K-Means Mood Clustering")
plt.xlabel("Energy")
plt.ylabel("Valence")
plt.grid(True)
plt.tight_layout()
plt.show()

```



```

[16] # ----- Apriori Algorithm -----
# Binarize continuous features
spotify_bin = pd.DataFrame()
for col in ['energy', 'danceability', 'valence']:
    spotify_bin[col + '_high'] = spotify_df[col] > spotify_df[col].median()

frequent_items = apriori(spotify_bin, min_support=0.1, use_colnames=True)
rules = association_rules(frequent_items, metric="lift", min_threshold=1.0)
print("Apriori Rules:\n", rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']].head())

```


Apriori Rules:

	antecedents	consequents	support	confidence	lift
0	(energy_high)	(danceability_high)	0.270843	0.541912	1.085181
1	(danceability_high)	(energy_high)	0.270843	0.542364	1.085181
2	(valence_high)	(energy_high)	0.318449	0.638365	1.277261
3	(energy_high)	(valence_high)	0.318449	0.637164	1.277261
4	(valence_high)	(danceability_high)	0.330559	0.662640	1.326938


```

#correlation heatmap for audio features
plt.figure(figsize=(12,8))
audio_features = ['acousticness', 'danceability', 'energy', 'instrumentalness',
                  'liveness', 'loudness', 'speechiness', 'tempo', 'valence']
corr = spotify_df[audio_features].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Between Audio Features")
plt.show()

```

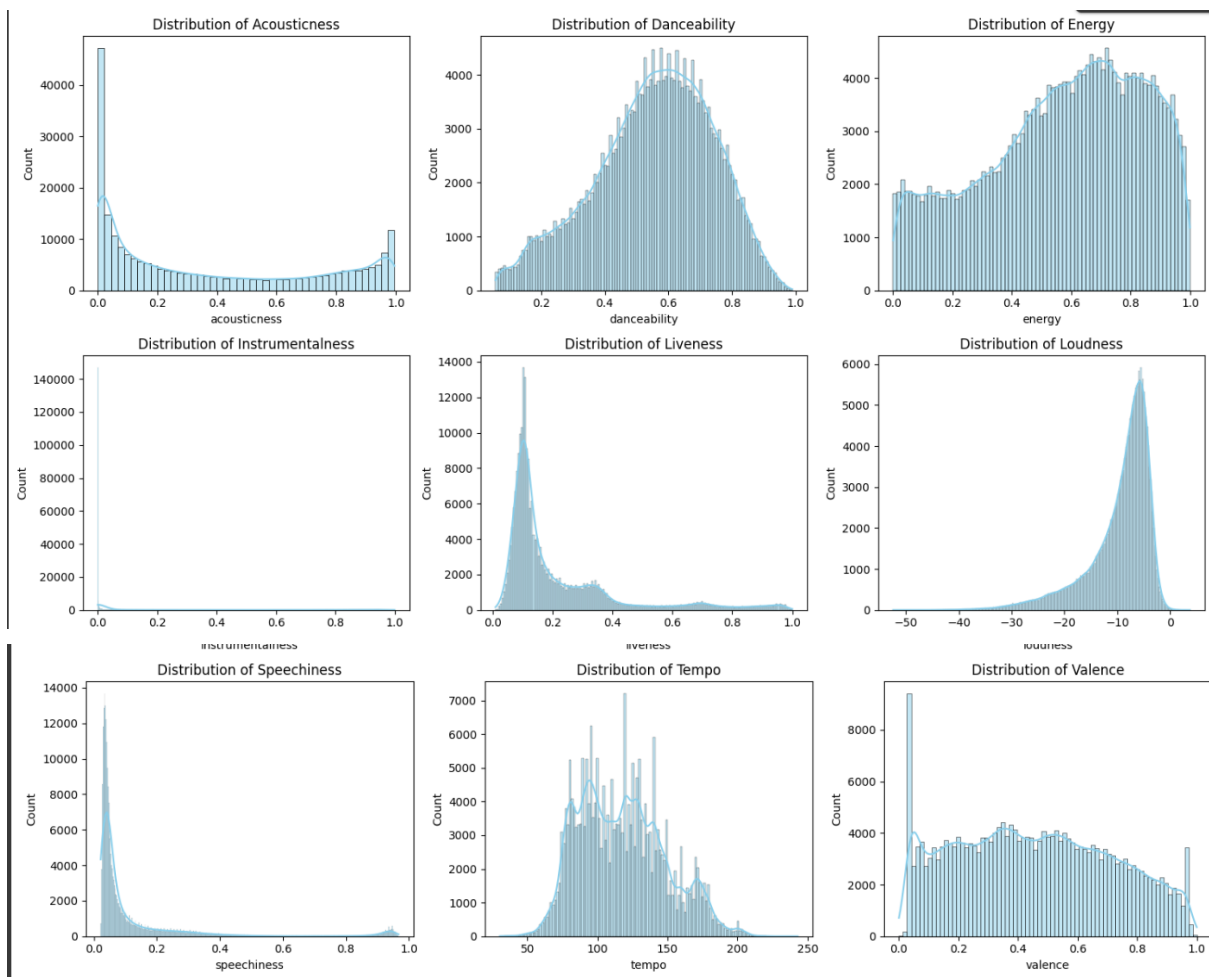


```

# Distribution Plot for All Audio Features
features_to_plot = ['acousticness', 'danceability', 'energy', 'instrumentalness',
                    'liveness', 'loudness', 'speechiness', 'tempo', 'valence']

plt.figure(figsize=(15,12))
for i, feature in enumerate(features_to_plot):
    plt.subplot(3, 3, i+1)
    sns.histplot(spotify_df[feature], kde=True, color='skyblue')
    plt.title(f"Distribution of {feature.capitalize()}")
plt.tight_layout()
plt.show()

```




```
[41] #Top Artists by Popularity
top_artists = spotify_df.groupby('artist_name')['popularity'].mean().sort_values(ascending=False).head(10)
plt.figure(figsize=(10,5))
sns.barplot(x=top_artists.values, y=top_artists.index, palette='rocket')
plt.title("Top 10 Artists by Average Popularity")
plt.xlabel("Average Popularity")
plt.ylabel("Artist")
plt.tight_layout()
plt.show()
```

`<ipython-input-41-137e4a242a36>:4: FutureWarning:`

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'y' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(x=top_artists.values, y=top_artists.index, palette='rocket')
```



Thank
you

The text "Thank you" is written in a dark blue, elegant cursive script. It is surrounded by a decorative arrangement of stylized flowers, leaves, and small dots. There are three yellow flowers with blue centers and two red flowers with green centers. The leaves are green with yellow veins. Small dots in blue, yellow, and red are scattered around the text, creating a festive and artistic feel.

